

Abstract :

The main aim is to Create a classification model to predict whether a person makes over \$50k a year

Data Understanding :

```
In [1]: import pandas as pd
a=['Age', 'Workclass', 'Fnlwgt', 'Education', 'education_num', 'marital_status', 'occupation', 'relationship', 'race', 'sex', 'capital_gain', 'capital_loss', '
data=pd.read_csv("C:\Users\zeus\Downloads\adult.csv",name=a)
data:

Out[1]:
```

	Age	Workclass	Fnlwgt	Education	education_num	marital_status	occupation	relationship	race	sex	capital_gain	capital_loss	hours_per_week	native_country	Income
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Nat-in-family	White	Male	2174	0	40	United-States	<=50K
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male	0	0	13	United-States	<=50K
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Nat-in-family	White	Male	0	0	40	United-States	<=50K
3	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	0	0	40	United-States	<=50K
4	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female	0	0	40	Cuba	<=50K
...
32556	27	Private	257302	Assoc-acdm	12	Married-civ-spouse	Tech-support	Wife	White	Female	0	0	38	United-States	<=50K
32557	40	Private	154374	HS-grad	9	Married-civ-spouse	Machine-op-inspct	Husband	White	Male	0	0	40	United-States	>50K
32558	58	Private	151910	HS-grad	9	Widowed	Adm-clerical	Unmarried	White	Female	0	0	40	United-States	<=50K
32559	22	Private	201490	HS-grad	9	Never-married	Adm-clerical	Own-child	White	Male	0	0	20	United-States	<=50K
32560	52	Self-emp-inc	287927	HS-grad	9	Married-civ-spouse	Exec-managerial	Wife	White	Female	15024	0	40	United-States	>50K

32561 rows x 15 columns

```
In [2]: data.isnull().sum() #Checking if there are any null values in data set

Out[2]: Age 0
Workclass 0
Fnlwgt 0
Education 0
education_num 0
marital_status 0
occupation 0
relationship 0
race 0
sex 0
capital_gain 0
capital_loss 0
hours_per_week 0
native_country 0
income: int64
dtype: int64

Data Cleaning :
```

```
In [3]: # Checking if there are any unwanted values in data set

(data['Age']== '?').sum()
(data['Workclass']== '?').sum()
(data['Education']== '?').sum()
(data['education_num']== '?').sum()
(data['marital_status']== '?').sum()
(data['occupation']== '?').sum()
(data['relationship']== '?').sum()
(data['race']== '?').sum()
(data['sex']== '?').sum()
(data['capital_gain']== '?').sum()
(data['capital_loss']== '?').sum()
(data['hours_per_week']== '?').sum()
(data['native_country']== '?').sum()
(data['income']== '?').sum()

Out[3]: 0

Now checking count of values in each row to replace unwanted value with this :
```

```
In [4]: data['Workclass'].value_counts()

Out[4]: Private 22696
Self-emp-not-inc 2541
Local-gov 2093
7 1836
State-gov 1298
Self-emp-inc 1116
Federal-gov 960
Without-pay 14
Never-worked 7
Name: Workclass, dtype: int64
```

```
In [5]: data['occupation'].value_counts()

Out[5]: Prof-specialty 4140
Craft-repair 4099
Exec-managerial 4066
Adm-clerical 3770
Sales 3650
Other-service 3293
Machine-op-inspct 2002
7 1843
Transport-moving 1597
Handlers-cleaners 1370
Farming-fishing 994
Tech-support 928
Protective-serv 649
Priv-house-serv 149
Armed-Forces 9
Name: occupation, dtype: int64
```

```
In [6]: data['native_country'].value_counts()

Out[6]: United-States 29170
Mexico 643
7 583
Philippines 198
Germany 137
Canada 121
Puerto-Rico 114
El-Salvador 106
India 100
Cuba 95
England 90
Jamaica 81
South 80
China 75
Italy 73
Dominican-Republic 70
Vietnam 67
Guatemala 64
Japan 62
Poland 60
Columbia 59
Taiwan 51
Haiti 44
Iran 43
Portugal 37
Nicaragua 34
Peru 31
France 29
Greece 29
Ecuador 28
Ireland 24
Hong 20
TrinidadTobago 19
Cambodia 19
Thailand 18
Laos 18
Yugoslavia 16
Outlying-US(Guam-USVI-etc) 14
Honduras 13
Hungary 13
Scotland 12
Noland-Netherlands 1
Name: native_country, dtype: int64
```

```
In [7]: #Replacing the unwanted values
data['Workclass']=data['Workclass'].replace(to_replace='?',value='Private')
data['occupation']=data['occupation'].replace(to_replace='?',value='Prof-specialty')
data['occupation']=data['occupation'].replace(to_replace='?',value='Prof-specialty')
data['native_country']=data['native_country'].replace(to_replace='?',value='United-States')
```

```
In [8]: data.columns

Out[8]: Index(['Age', 'Workclass', 'Fnlwgt', 'Education', 'education_num', 'marital_status', 'occupation', 'relationship', 'race', 'sex', 'capital_gain', 'capital_loss', 'hours_per_week', 'native_country', 'income'],
dtype='object')
```

As all the values should be in integer datatype , string values has to be converted into integer values by using LabelEncoder

```
In [9]: from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()
data['Workclass']=le.fit_transform(data['Workclass'])
data['Education']=le.fit_transform(data['Education'])
data['marital_status']=le.fit_transform(data['marital_status'])
data['occupation']=le.fit_transform(data['occupation'])
data['relationship']=le.fit_transform(data['relationship'])
data['race']=le.fit_transform(data['race'])
data['sex']=le.fit_transform(data['sex'])
data['race']=le.fit_transform(data['sex'])
data['native_country']=le.fit_transform(data['native_country'])
```

```
In [10]: data.shape

Out[10]: (32561, 15)
```

Splitting dataset into train and test data :

```
In [11]: x=data.iloc[:,1:14].values
y=data.iloc[:,14].values
x.shape

Out[11]: (32561, 14)

In [12]: y.shape

Out[12]: (32561,)
```

Data Preprocessing :

```
In [13]: from sklearn.model_selection import train_test_split
xtrain,xtest,ytrain,ytest=train_test_split(x,y,test_size=0.20,random_state=1)

In [14]: xtrain.shape

Out[14]: (26048, 14)

In [15]: ytrain.shape

Out[15]: (26048,)
```

Model Building :

```
In [16]: from sklearn.metrics import confusion_matrix,classification_report
from sklearn.metrics import accuracy_score

In [17]: from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
```

```
In [18]: #creating function for all models
def apply_model(xtrain,xtest,ytrain,ytest,model):
    model.fit(xtrain,ytrain)
    ypred=model.predict(xtest)
    print("prediction of income for test data :", ypred)
    print("Training score:",model.score(xtrain,ytrain))
    print("Testing score:",model.score(xtest,ytest))
    print("Accuracy score:",accuracy_score(ytest,ypred)*100)
    cm=confusion_matrix(ytest,ypred)
    print("Confusion matrix: \n",cm)
    print("Classification report : \n ",classification_report(ytest,ypred))
    pr=cm[0][0]/(cm[0][0]+cm[1][0])
    se=cm[0][1]/(cm[0][0]+cm[0][1])
    print("recall:",pr)
    print("F-i score: ",2*pr*se/(pr+se))
    acc=(cm[0][0]+cm[1][1])/(cm[0][0]+cm[1][0]+cm[0][1]+cm[1][1])
    print("Accuracy: ",acc)
    print(" ")
    mis=(cm[1][0]+cm[0][1])/(cm[0][0]+cm[1][0]+cm[0][1]+cm[1][1])
    print("percentage of misclassification : \n",mis)
```

Model 1: Decision Tree Classifier

```
In [19]: Model1=DecisionTreeClassifier(criterion='gini')
apply_model(xtrain,xtest,ytrain,ytest,Model1)

prediction of income for test data : [' <=50K' ' <=50K' ' >50K' ... ' <=50K' ' <=50K' ' >50K']
Training score: 0.9999616093366094
Testing score: 0.8129894058037771
Accuracy score: 81.2989405803777
confusion matrix:
[[4365 661]
 [ 357  801]]
classification report :
              precision    recall  f1-score   support

<=50K      0.89      0.87      0.88      5026
>50K       0.58      0.63      0.60      1487

accuracy      0.81      0.75      0.74      6513
macro avg     0.74      0.75      0.74      6513
weighted avg   0.82      0.81      0.82      6513

precision: 0.886834620073141
recall: 0.868483883804218
F-1 score: 0.877563293124246
Accuracy: 0.8129894058037771

percentage of misclassification :
57.101489328043
```

Model 2: Random Forest Classifier

```
In [20]: Model2=RandomForestClassifier()
apply_model(xtrain,xtest,ytrain,ytest,Model2)

prediction of income for test data : [' <=50K' ' <=50K' ' >50K' ... ' <=50K' ' <=50K' ' >50K']
Training score: 0.9999616093366094
Testing score: 0.8604329801934593
Accuracy score: 86.04329801934593
confusion matrix:
[[4647 379]
 [ 530  871]]
classification report :
              precision    recall  f1-score   support

<=50K      0.90      0.92      0.91      5026
>50K       0.72      0.64      0.68      1487

accuracy      0.81      0.78      0.79      6513
macro avg     0.81      0.78      0.79      6513
weighted avg   0.86      0.86      0.86      6513

precision: 0.8976241066254588
recall: 0.924592120970951
F-1 score: 0.9109085303069685
Accuracy: 0.8604329801934593

percentage of misclassification :
530.0581913096884
```

Model 3: Logistic Regression

```
In [21]: Model3=LogisticRegression(solver='liblinear')
apply_model(xtrain,xtest,ytrain,ytest,Model3)

prediction of income for test data : [' <=50K' ' <=50K' ' >50K' ... ' <=50K' ' <=50K' ' >50K']
Training score: 0.7916338697786998
Testing score: 0.8020881314294488
Accuracy score: 80.20881314294488
confusion matrix:
[[4792 234]
 [1055 432]]
classification report :
              precision    recall  f1-score   support

<=50K      0.82      0.95      0.88      5026
>50K       0.65      0.29      0.40      1487

accuracy      0.73      0.62      0.64      6513
macro avg     0.78      0.80      0.74      6513
weighted avg   0.76      0.77      0.76      6513

precision: 0.8195655891910382
recall: 0.933442101744131
F-1 score: 0.8814494619700175
Accuracy: 0.8020881314294488

percentage of misclassification :
1055.0359281437127
```

Model 4: KNN Classifier

```
In [22]: Model4=KNeighborsClassifier(n_neighbors=3)
apply_model(xtrain,xtest,ytrain,ytest,Model4)

prediction of income for test data : [' <=50K' ' <=50K' ' >50K' ... ' <=50K' ' <=50K' ' >50K']
Training score: 0.8624462350712531
Testing score: 0.7729157070474436
Accuracy score: 77.29157070474436
confusion matrix:
[[4440 586]
 [ 893 594]]
classification report :
              precision    recall  f1-score   support

<=50K      0.83      0.88      0.86      5026
>50K       0.50      0.40      0.45      1487

accuracy      0.67      0.64      0.77      6513
macro avg     0.67      0.64      0.65      6513
weighted avg   0.76      0.77      0.76      6513

precision: 0.8325520345021564
recall: 0.813462897060088
F-1 score: 0.8572256009267305
Accuracy: 0.7729157070474436

percentage of misclassification :
893.0899738983571
```

Model 5: SVC Classifier (with Linear Kernel)

```
In [23]: Model5=SVC(kernel='linear',C=1)
apply_model(xtrain,xtest,ytrain,ytest,Model5)

prediction of income for test data : [' <=50K' ' <=50K' ' >50K' ... ' <=50K' ' <=50K' ' >50K']
Training score: 0.7929207616707616
Testing score: 0.8053124520190389
Accuracy score: 80.53124520190389
confusion matrix:
[[5018  8]
 [1260 227]]
classification report :
              precision    recall  f1-score   support

<=50K      0.80      1.00      0.89      5026
>50K       0.97      0.15      0.26      1487

accuracy      0.89      0.58      0.81      6513
macro avg     0.88      0.58      0.58      6513
weighted avg   0.84      0.81      0.75      6513

precision: 0.7992991398534565
recall: 0.9984082769598089
F-1 score: 0.8878273177636234
Accuracy: 0.8053124520190389

percentage of misclassification :
1260.0012283126055
```

Conclusion:

Random Forest is the best model in terms of accuracy